



# Disease Prediction Using Big Data Analytics with PySpark

Intro to Big Data Analytics, Fall 2025

TEAM 3

*Laman Panakhova, Shahla Azizova, Khadija Ahmadova,  
Mahbuba Jafarzada, Emil Hajiyev*

# The Need: Healthcare Is Becoming a Data-Heavy Industry

Hospitals and medical systems now generate enormous amounts of data:

- Electronic health records
- Symptom logs
- Lab reports
- Wearable device streams

This data **can** help identify disease risk early.

But only if we can process it efficiently and consistently.

Traditional tools like Pandas or Excel cannot handle this growth.

Healthcare analytics needs a **scalable, distributed system**.

# The Problem: High-Dimensional Data That Doesn't Fit Traditional Tools

We worked with a healthcare dataset that already shows this challenge:

- 4920 patient samples + 1230 test samples
- 41 symptoms per patient → ***wide*** dataset
- Multi-disease classification
- Strict need for consistency and reliability

Even though this dataset is small, real medical datasets are much larger and cannot be processed on a single machine.

So the real question becomes:

**How do we design a solution that can scale to real clinical workloads?**

# The Data

Train rows: 4920 columns: 134  
Test rows: 42 columns: 133

	itching	skin_rash	nodal_skin_eruptions	continuous_sneezing	shivering	chills	joint_pain	stomach_pain	acidity	ulcers_on_tongue	...	scurring	skin_peeling	
0	1	1		1	0	0	0	0	0	0	0	...	0	0
1	0	1		1	0	0	0	0	0	0	0	...	0	0
2	1	0		1	0	0	0	0	0	0	0	...	0	0
3	1	1		0	0	0	0	0	0	0	0	...	0	0
4	1	1		1	0	0	0	0	0	0	0	...	0	0

<u>silver_like_dusting</u>	<u>small_dents_in_nails</u>	<u>inflammatory_nails</u>	<u>blister</u>	<u>red_sore_around_nose</u>	<u>yellow_crust_ooze</u>	<u>prognosis</u>	<u>_c133</u>
0	0	0	0	0	0	Fungal infection	None
0	0	0	0	0	0	Fungal infection	None
0	0	0	0	0	0	Fungal infection	None
0	0	0	0	0	0	Fungal infection	None
0	0	0	0	0	0	Fungal infection	None

# Why Spark?

Spark addresses the bottlenecks that traditional tools face:

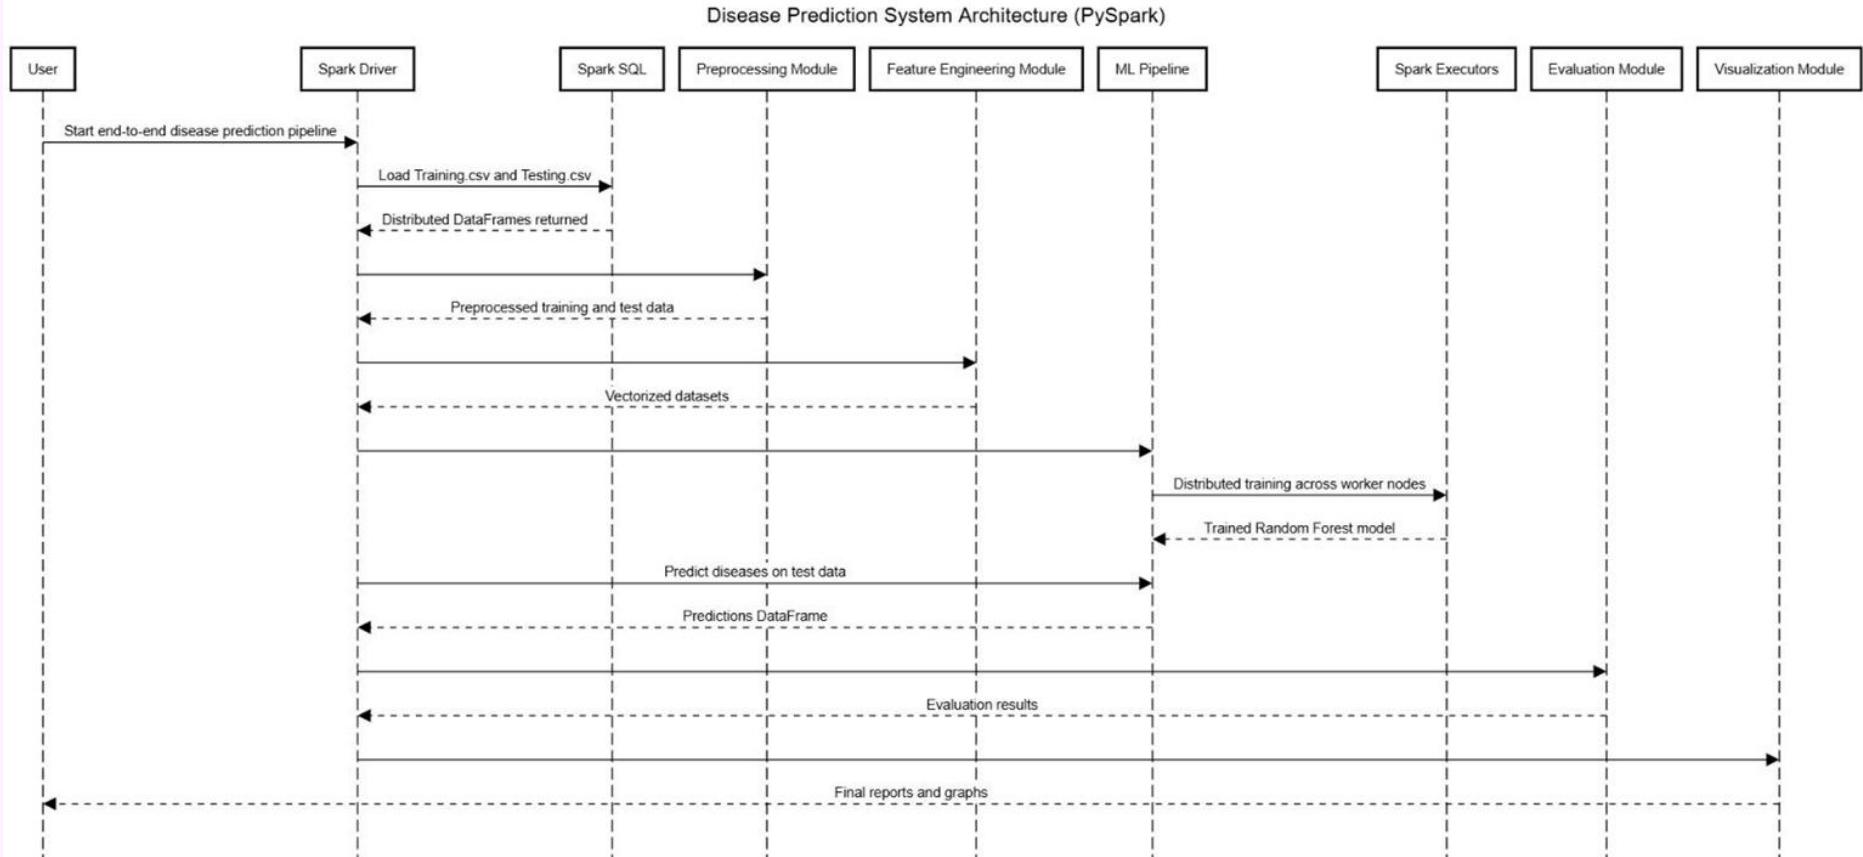
- Distributed processing across multiple nodes
- In-memory computation for speed
- Fault tolerance through RDD lineage
- Optimized SQL engine (Catalyst)
- MLLib for distributed machine learning

Instead of struggling with RAM limits, Spark **divides work across a cluster.**

Perfect for healthcare-scale datasets.



# Our Approach: Build a Big-Data Pipeline for Disease Prediction



# How We Prepared the Data (Distributed Preprocessing)

We cleaned and prepared the dataset using distributed DataFrame transformations:

- Normalized inconsistent column names
- Converted 41 symptom features into numeric format
- Applied label encoding (StringIndexer)
- Created vectorized feature columns

Each transformation ran **in parallel**, not sequentially.  
This demonstrates how Spark handles even wide, complex datasets.



# Distributed EDA: Understanding the Data at Scale

Instead of analyzing symptoms on a single machine, Spark allowed us to compute:

- Symptom frequency
- Disease distribution
- Co-occurrence patterns

These operations used Spark SQL and DataFrames, meaning they could scale to millions of records with very little change.

The key insight: **EDA becomes scalable, repeatable, and efficient.**

# The Solution in Action: Spark MLlib Pipeline & DEMO

**VectorAssembler**  
packs 41 features  
in parallel

**RandomForestClassifier:**  
trains multiple trees  
simultaneously

**CrossValidator:**  
distributes  
evaluation across  
executors

The pipeline is not only for prediction—  
it demonstrates **how Spark enables large-scale analytical workflows.**

## Results and What They Tell Us

Even on this dataset, the distributed workflow was effective:

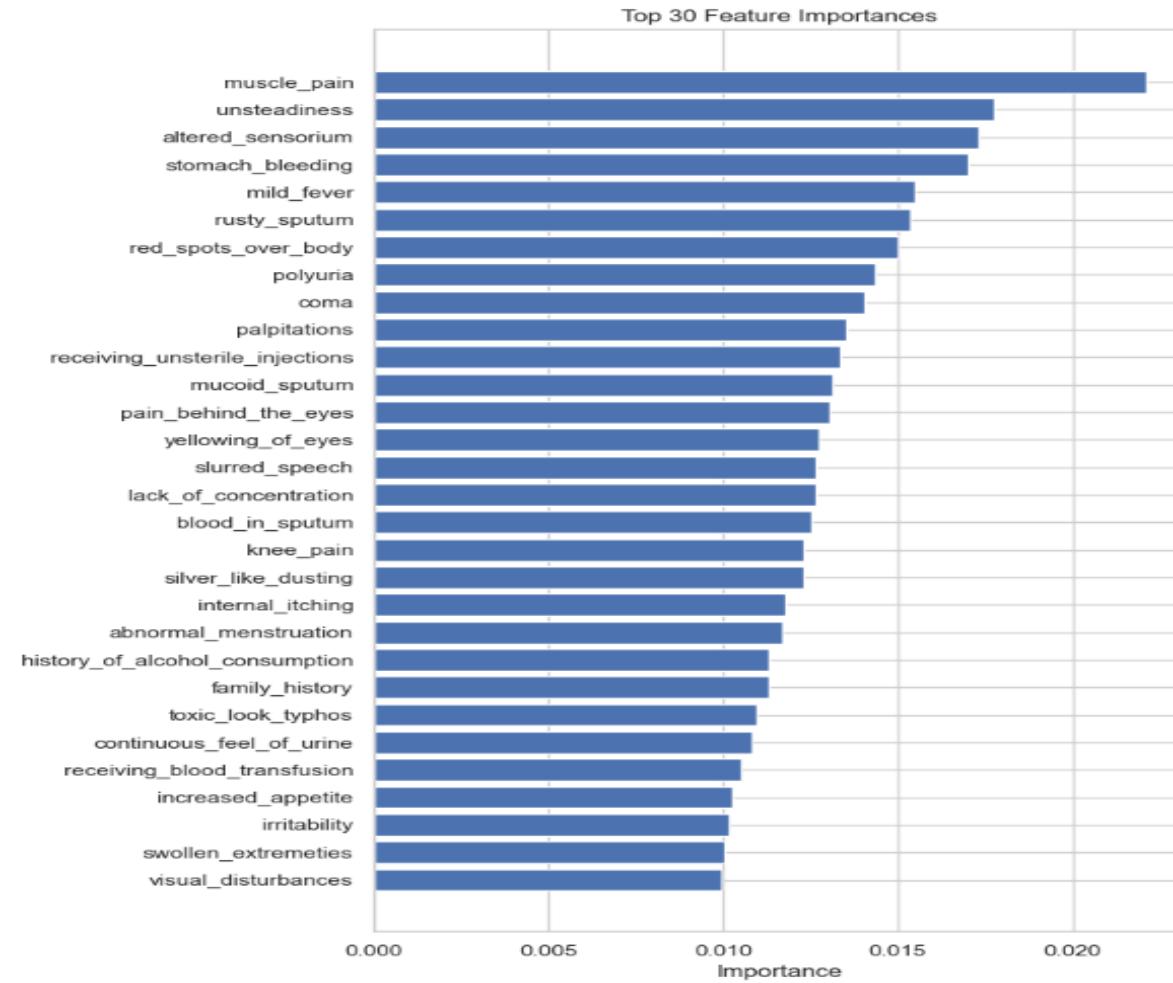
- Accuracy: **98%**
- Precision, recall, and F1 all above **0.98**

But the real point is not the number.

It is that **the pipeline works, and it scales.**

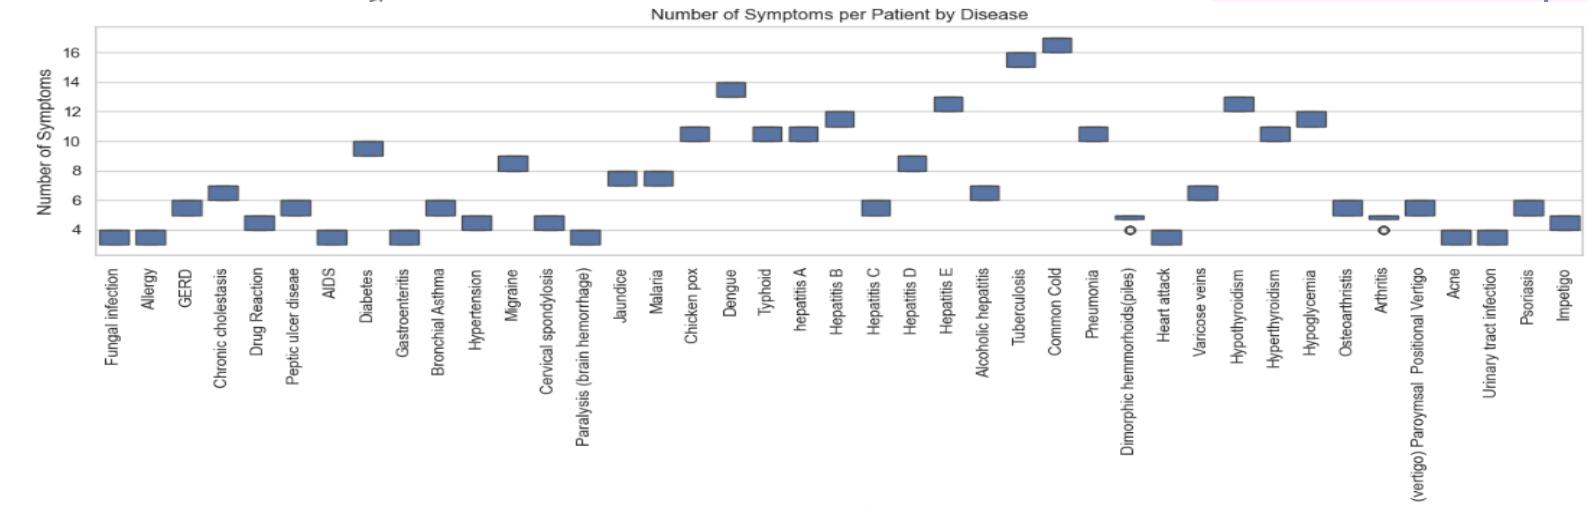
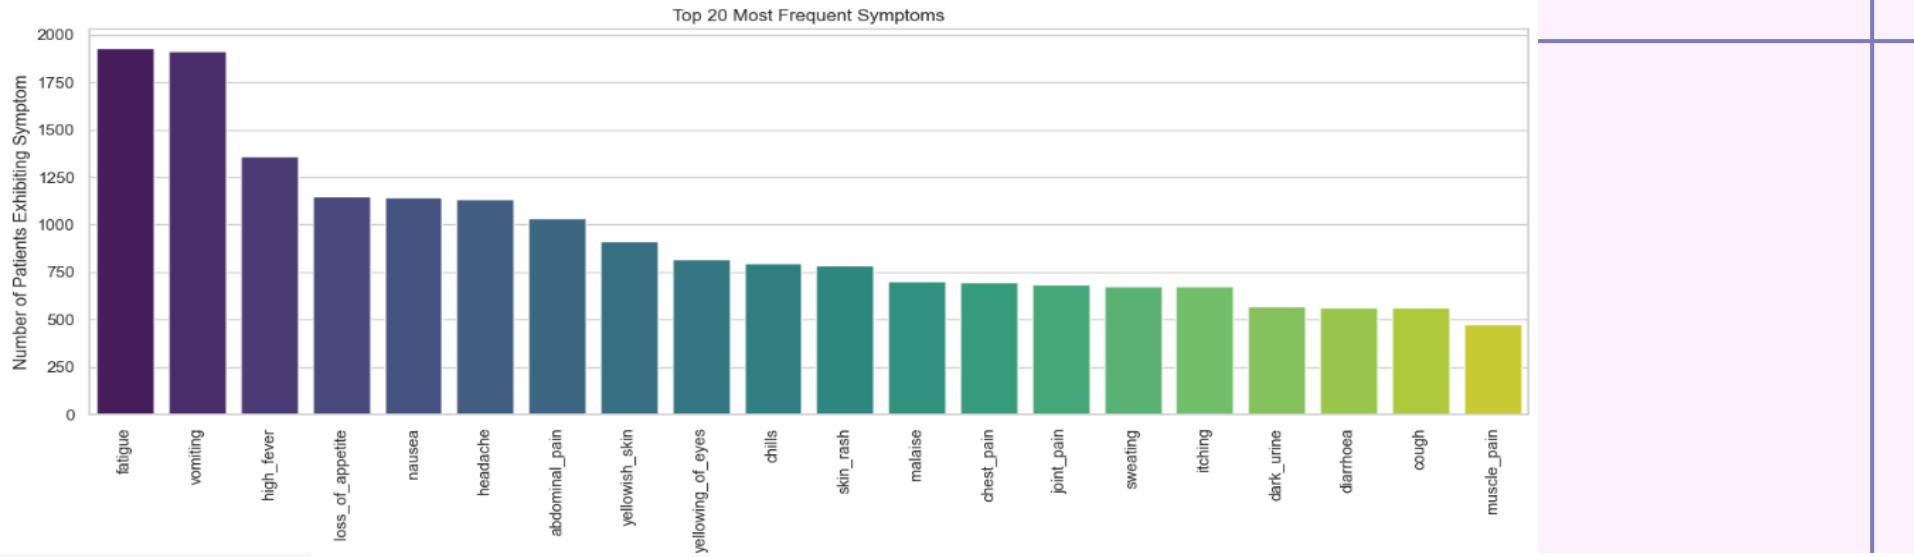
If tomorrow we receive 10 million patient records, Spark can run the same workflow with minimal changes.





[19]:

	feature	importance
0	muscle_pain	0.022095
1	unsteadiness	0.017738
2	altered_sensorium	0.017337
3	stomach_bleeding	0.017010
4	mild_fever	0.015493
5	rusty_sputum	0.015344
6	red_spots_over_body	0.015019
7	polyuria	0.014352
8	coma	0.014048
9	palpitations	0.013516
10	receiving_unsterile_injections	0.013335
11	mucoid_sputum	0.013122
12	pain_behind_the_eyes	0.013044
13	yellowing_of_eyes	0.012763
14	slurred_speech	0.012657
15	lack_of_concentration	0.012636
16	blood_in_sputum	0.012531
17	knee_pain	0.012307
18	silver_like_dusting	0.012295
19	internal_itching	0.011768
20	abnormal_menstruation	0.011710
21	history_of_alcohol_consumption	0.011312
22	family_history	0.011285
23	toxic_look_typhos	0.010966
24	continuous_feel_of_urine	0.010821
25	receiving_blood_transfusion	0.010509
26	increased_appetite	0.010251
27	irritability	0.010179
28	swollen_extremeties	0.010043
29	visual_disturbances	0.009971



# Challenges, Lessons, and Big-Data Impact

## Challenges we faced

- Inconsistent column names across CSVs
- Wide, 41-feature dataset
- Label mismatches between training/testing
- Spark's limited visualization environment

## Big-Data Impact

- Pipeline demonstrates readiness for large, real clinical datasets
- Spark can support early disease prediction at scale
- Provides foundation for hospital-level analytics systems



# Conclusion & Future Directions

## Conclusion

Our project shows that:

- Healthcare needs scalable analytical pipelines
- Spark provides distributed, fault-tolerant infrastructure
- A well-designed data pipeline is essential for real-world disease prediction
- The focus is not only on prediction—but on **scalable, reliable data processing**

## Future Work

- Move to HDFS or cloud object storage
- Integrate streaming data (Spark Structured Streaming)
- Scale to millions of patient records
- Add more advanced distributed models (GBTs, neural networks)
- Deploy as a real-time decision-support system for hospitals



Thank you!

