

School of Information Technologies and Engineering, ADA University

INFT4836 – Intro to Big Data Analytics

Fall 2025 – 11/29/2025

Assignment 2 by Laman Panakhova 16882

Apache Spark

Instructions:

Overview

In this assignment, you will apply the concepts learned in the course by working with **Apache Spark**, including **RDDs**, **Spark SQL**, **DataFrames**, and an introduction to **basic data preprocessing / feature engineering** using Spark. Your goal is to work with a real-world large dataset and demonstrate your understanding of distributed data processing, querying, and transformation.

You will complete **three major tasks** using the **same dataset** throughout the assignment. These tasks will allow you to practice data manipulation with RDDs, build SQL queries over distributed datasets, and perform more advanced operations using DataFrames.

Please submit a **single PDF file** containing all answers, code snippets, outputs (screenshots), figures, and explanations.

Assignment Tasks

Task 1: Spark RDDs – Core Transformations & Actions (40 points)

Dataset Selection (no points, required)

Choose a dataset from a domain of your choice (healthcare, finance, IoT, cybersecurity, retail, sports analytics, transportation, or others). Your dataset should:

Be at least **100,000 rows** or reasonably large

Contain **multiple features** (numeric + categorical preferred)

Be obtained from a reliable source (Kaggle, UCI ML Repository, Open Data Portal, etc.)

Setup (5 points)

Install and configure Apache Spark on your local machine.

Load your dataset using standard Spark methods (`textFile`, `wholeTextFiles`, or `csv reader`).

Provide screenshots verifying successful installation and dataset loading.

RDD Transformations & Actions (35 points)

Using Spark RDDs, perform the following operations **with explanations of why each step is meaningful for your dataset**:

1. **map()** transformation – derive a new feature or convert raw entries into structured tuples (7 pts)
2. **filter()** transformation – remove invalid / missing / noisy records (7 pts)
3. **flatMap()** transformation – break complex fields (e.g., text, lists, logs) into smaller elements (7 pts)
4. **reduceByKey() or aggregate()** – compute summary statistics or grouped aggregates (7 pts)
5. **take() or collect() (action)** – extract sample outputs and comment on distributed execution (7 pts)

Include relevant code, outputs, and short justifications for each transformation.

Task 2: Spark SQL – Querying Distributed Data (30 points)

*Load the same dataset using Spark SQL and create a **temporary view**.*

Using Spark SQL, perform and document the following:

1. Aggregation Query (10 pts) *Example: average value by category, count by class label, sum by month, etc.*
2. Filtering Query (10 pts) *Use a WHERE clause with logical operators or range conditions.*
3. Join Query (10 pts) - *Create (or simulate) a second small helper table (e.g., category descriptions, lookup table) and perform an INNER or LEFT JOIN. Explain why the join is meaningful for analysis. Include SQL query text, results, and screenshots/output tables.*

Task 3: Spark DataFrames – Data Cleaning & Transformations (30 points)

Use Spark DataFrames to explore and preprocess your dataset.

Perform the following DataFrame operations:

1. Data Cleaning: Handling Missing or Incorrect Values (10 pts) - *Demonstrate .dropna(), .fillna(), or custom cleaning logic. Explain why the chosen strategy is appropriate.*
2. groupBy() + agg() operations (10 pts) - *Compute multiple aggregate functions (e.g., mean, stddev, count) on at least one categorical field.*
3. orderBy() / sort() with multiple columns (10 pts) - *Show sorting results and comment on distributed execution.*

Include code, results, and interpretations for each step.

Solutions:

In this assignment we focus on using Apache Spark to process the *Online Retail II* dataset using three different APIs: RDDs, Spark SQL, and DataFrames. The goal is to demonstrate understanding of data loading, cleaning, transformation, querying, and performing basic analytical tasks at scale. All tasks were completed in a local PySpark environment.

Description of the dataset:

The dataset used in this project is the Online Retail II dataset, containing real transaction records from a UK-based online retail company. Each row represents a single product purchased within an invoice, with details such as *Invoice ID*, *StockCode*, *Description*, *Quantity*, *InvoiceDate*, *Price*, *Customer ID*, and *Country*. This dataset is commonly used in data analytics and e-commerce research to study customer behavior, analyze sales trends, forecast demand, and practice data cleaning and processing with real-world transactional data. Its structure allows exploration of revenue patterns, customer segmentation, and other insights relevant to online retail operations.

You can access the dataset here: <https://www.kaggle.com/datasets/mashlyn/online-retail-ii-uci>

The main criteria for choosing the dataset were below per assignment instruction:

- ✓ Be at least 100,000 rows or reasonably large,
- ✓ Contain multiple features (numeric + categorical preferred),
- ✓ Be obtained from a reliable source (Kaggle, UCI ML Repository, Open Data Portal, etc.)

Task 1: Spark RDD - Core Transformations & Actions

To begin the analysis, Apache Spark was installed and configured on the local machine to provide a scalable environment for handling large datasets. The dataset was loaded using Spark's standard CSV reader, which efficiently reads structured data into a DataFrame for processing. Spark's methods allow filtering, transformation, and aggregation on the dataset with high performance, even on modest hardware. Screenshots of the Spark session and the loaded dataset verify successful installation and loading, demonstrating that the environment is ready for further data analysis.

The screenshot shows a Jupyter Notebook interface. The top cell contains Python code for installing PySpark via pip:

```
!pip install pyspark
```

Output shows requirements already satisfied:

```
Requirement already satisfied: pyspark in /usr/local/lib/python3.12/dist-packages (3.5.1)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.12/dist-packages (from pyspark) (0.10.9.7)
```

The second cell contains code to create a SparkContext and print its first three elements:

```
!pip install pyspark
from pyspark import SparkContext
sc = SparkContext()

print(sc.parallelize([1,2,3]).take(3))
```

Output shows requirements already satisfied and the result [1, 2, 3]:

```
Requirement already satisfied: pyspark in /usr/local/lib/python3.12/dist-packages (3.5.1)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.12/dist-packages (from pyspark) (0.10.9.7)
[1, 2, 3]
```

The third cell contains code to print the PySpark version:

```
import pyspark
print(pyspark.__version__)
```

Output shows the version 3.5.1:

```
3.5.1
```

The bottom cell contains a longer PySpark session setup code:

```
# Imports and version check
import pyspark
print(pyspark.__version__)

# Cell: common setup
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, to_timestamp, month, year, when, sum as spark_sum, avg, stddev, count
from pyspark.sql.types import StructType, StructField, IntegerType, StringType, DoubleType, TimestampType

# create spark session
spark = SparkSession.builder \
    .appName("NFT4836_Hw2_OnlineRetailIIiI") \
    .master("local[1*]") \
    .config("spark.driver.memory", "6G") \
    .getOrCreate()

sc = spark.sparkContext

print("Spark version:", spark.version)
```

Output shows the version 3.5.1:

```
3.5.1
Spark version: 3.5.1
```

Furthermore, the dataset was loaded, and several required transformations were implemented using both RDD operations and their equivalent DataFrame APIs.

✓ *Dataset Loading*

The CSV file was read with header and schema inference enabled. A sample of the dataset was inspected using `.show()` to confirm successful loading.

```

# Cell: load CSV as DataFrame (with inferSchema disabled initially then explicit schema)
file_path = r"/content/online_retail_II.csv"

# If CSV has header and default separators, read with Spark CSV reader:
df_raw = spark.read.option("header", True).option("inferSchema", False).option("multiLine", True).option("escape", "\\'")
print("Columns:", df_raw.columns)
print("Count (rows):", df_raw.count())
df_raw.show(5, truncate=False)

Columns: ['Invoice', 'StockCode', 'Description', 'Quantity', 'InvoiceDate', 'Price', 'Customer ID', 'Country']
Count (rows): 259454
+-----+-----+-----+-----+-----+-----+-----+
|Invoice|StockCode|Description|Quantity|InvoiceDate|Price|Customer ID|Country|
+-----+-----+-----+-----+-----+-----+-----+
|489434 |85048   |15CM CHRISTMAS GLASS BALL 20 LIGHTS|12    |2009-12-01 07:45:00|6.95 |13085.0    |United Kingdom|
|489434 |79323P  |PINK CHERRY LIGHTS                |12    |2009-12-01 07:45:00|6.75 |13085.0    |United Kingdom|
|489434 |79323W  |WHITE CHERRY LIGHTS                |12    |2009-12-01 07:45:00|6.75 |13085.0    |United Kingdom|
|489434 |22041   |RECORD FRAME 7" SINGLE SIZE        |48    |2009-12-01 07:45:00|2.1  |13085.0    |United Kingdom|
|489434 |21232   |STRAWBERRY CERAMIC TRINKET BOX    |24    |2009-12-01 07:45:00|1.25 |13085.0    |United Kingdom|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```

After doing some preprocessing steps on data let's continue with below transformations:

- ✓ *map() transformation – derive a new feature or convert raw entries into structured tuples*

A new column TotalValue = Quantity × UnitPrice was created.

- RDD version: by mapping each row tuple
- DF version: using withColumn() and col()

This identifies revenue per invoice line.

```

# Now create an RDD of tuples
rdd = df.rdd.map(lambda row: (
    row['Invoice'],
    row['StockCode'],
    row['Description'],
    row['Quantity'] if row['Quantity'] is not None else 0,
    row['InvoiceDateTS'],
    row['Price'] if row['Price'] is not None else 0.0,
    row['Customer ID'],
    row['Country']
))
# sample
print("Sample RDD tuples:")
for t in rdd.take(5):
    print(t)
root
|-- Invoice: string (nullable = true)
|-- StockCode: string (nullable = true)
|-- Description: string (nullable = true)
|-- Quantity: integer (nullable = true)
|-- InvoiceDate: timestamp (nullable = true)
|-- Price: double (nullable = true)
|-- Customer ID: integer (nullable = true)
|-- Country: string (nullable = true)
|-- InvoiceDateTS: timestamp (nullable = true)

+-----+-----+-----+-----+-----+-----+-----+
|Invoice|StockCode|Description|Quantity|InvoiceDate|Price|Customer ID|Country|InvoiceDateTS|
+-----+-----+-----+-----+-----+-----+-----+
|489434 |85048   |15CM CHRISTMAS GLASS BALL 20 LIGHTS|12    |2009-12-01 07:45:00|6.95 |13085.0    |United Kingdom|2009-12-01 07:45:00|
|489434 |79323P  |PINK CHERRY LIGHTS                |12    |2009-12-01 07:45:00|6.75 |13085.0    |United Kingdom|2009-12-01 07:45:00|
|489434 |79323W  |WHITE CHERRY LIGHTS                |12    |2009-12-01 07:45:00|6.75 |13085.0    |United Kingdom|2009-12-01 07:45:00|
|489434 |22041   |RECORD FRAME 7" SINGLE SIZE        |48    |2009-12-01 07:45:00|2.1  |13085.0    |United Kingdom|2009-12-01 07:45:00|
|489434 |21232   |STRAWBERRY CERAMIC TRINKET BOX    |24    |2009-12-01 07:45:00|1.25 |13085.0    |United Kingdom|2009-12-01 07:45:00|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

Sample 800 tuples:
('489434', '85048', '15CM CHRISTMAS GLASS BALL 20 LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.95, 13085, 'United Kingdom')
('489434', '79323P', 'PINK CHERRY LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.75, 13085, 'United Kingdom')
('489434', '79323W', 'WHITE CHERRY LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.75, 13085, 'United Kingdom')
('489434', '22041', 'RECORD FRAME 7" SINGLE SIZE', 48, datetime.datetime(2009, 12, 1, 7, 45), 2.1, 13085, 'United Kingdom')
('489434', '21232', 'STRAWBERRY CERAMIC TRINKET BOX', 24, datetime.datetime(2009, 12, 1, 7, 45), 1.25, 13085, 'United Kingdom')

```

✓ *filter() transformation – remove invalid / missing / noisy records*

Rows with invalid values were removed:

- Negative or zero quantities
- Negative prices
- Null CustomerID entries

This ensures clean data for analysis.

```
# Filter invalid records
rdd_filtered = rdd.filter(lambda t: t[0] is not None and t[3] is not None and t[3] > 0 and t[5] is not None and t[5] > 0 and t[6] is not None and t[6] > 0)
print("Count before filter:", rdd.count())
print("Count after filter:", rdd_filtered.count())
# Show some samples
for t in rdd_filtered.take(5):
    print(t)

Count before filter: 850429
Count after filter: 805549
('489434', '85648', '15CM CHRISTMAS GLASS BALL 20 LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.95, 13085, 'United Kingdom')
('489434', '79323P', 'PINK CHERRY LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.75, 13085, 'United Kingdom')
('489434', '79323W', 'WHITE CHERRY LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.75, 13085, 'United Kingdom')
('489434', '22841', 'RECORD FRAME 7" SINGLE SIZE ', 48, datetime.datetime(2009, 12, 1, 7, 45), 2.1, 13085, 'United Kingdom')
('489434', '21232', 'STRAWBERRY CERAMIC TRINKET BOX', 24, datetime.datetime(2009, 12, 1, 7, 45), 1.25, 13085, 'United Kingdom')
```

✓ *flatMap() transformation – break complex fields (e.g., text, lists, logs) into smaller elements*

Product descriptions were tokenized into individual lowercase words.

This demonstrates how flatMap() can produce more output elements than input elements.

```
# flatMap: tokenizing descriptions into words (lowercase, simple split, remove empty)
desc_words = rdd_filtered.flatMap(lambda t: [(w.strip().lower()) for w in (t[2] or "").split() if w.strip() != ""])
# Count distinct words and top 20 frequent words
word_counts = desc_words.map(lambda w: (w, 1)).reduceByKey(lambda a, b: a + b)
top_words = word_counts.takeOrdered(20, key=lambda x: -x[1])
print("Top description words (word, count):")
for w, c in top_words:
    print(w, c)

Top description words (word, count):
of 70769
set 70506
bag 69889
red 69818
heart 59294
pink 45740
retrospot 45325
design 43086
vintage 42804
box 38404
white 38263
cake 35570
metal 35472
christmas 34466
blue 30160
hanging 28817
holder 27851
jumbo 27225
sign 26759
3 26420
```

✓ *reduceByKey() or aggregate() – compute summary statistics or grouped aggregates*

Two key-value aggregations were implemented:

- Total revenue per Country
- Total quantity per StockCode

This step shows how to aggregate distributed key-value data efficiently.

```
# revenue per row = Quantity * Price
rdd_revenue = rdd_filtered.map(lambda t: (t[7], float(t[3]) * float(t[5]) if float(t[3]) is not None and t[5] is not None else 0.0))
# sum revenue by country
revenue_by_country = rdd_revenue.reduceByKey(lambda a, b: a + b)
print("Top revenue by country (top 10):")
for c, tot in revenue_by_country.takeOrdered(10, key=lambda x: -x[1]):
    print(c, round(tot,2))

# total quantity by stock code
qty_by_stock = rdd_filtered.map(lambda t: (t[1], t[3])).reduceByKey(lambda a, b: a + b)
print("Top 10 stock codes by total quantity:")
for s, q in qty_by_stock.takeOrdered(10, key=lambda x: -x[1]):
    print(s, q)

Top revenue by country (top 10):
United Kingdom 14723147.52
EIRE 623631.11
Netherlands 554222.34
Germany 431262.46
France 355257.47
Australia 169968.11
Spain 109178.53
Switzerland 106365.34
Sweden 91549.72
Denmark 69862.19
Top 10 stock codes by total quantity:
84077 109169
850998 94983
85123A 93697
21212 91263
23843 80995
84879 79913
22197 77971
23166 77916
17003 71129
21977 55270
```

- ✓ *take() (action) – extract sample outputs and comment on distributed execution*

A small sample of the cleaned dataset (first 5 rows) was shown to verify transformations.

```
# take a small sample of filtered RDD
sample_rows = rdd_filtered.take(10)
print("Sample rows (first 10):")
for r in sample_rows:
    print(r)

Sample rows (first 10):
('489434', '85040', '15IN CHRISTMAS GLASS BALL 20 LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.95, 13085, 'United Kingdom')
('489434', '79323P', 'PINK CHERRY LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.75, 13085, 'United Kingdom')
('489434', '79323W', 'WHITE CHERRY LIGHTS', 12, datetime.datetime(2009, 12, 1, 7, 45), 6.75, 13085, 'United Kingdom')
('489434', '22841', 'RECORD FRAME 7" SINGLE SIZE', 48, datetime.datetime(2009, 12, 1, 7, 45), 2.1, 13085, 'United Kingdom')
('489434', '21232', 'STRAWBERRY CERAMIC TRINKET BOX', 24, datetime.datetime(2009, 12, 1, 7, 45), 1.25, 13085, 'United Kingdom')
('489434', '22664', 'PINK DOUGHNUT TRINKET POT', 24, datetime.datetime(2009, 12, 1, 7, 45), 1.65, 13085, 'United Kingdom')
('489434', '21871', 'SAVE THE PLANET MUG', 24, datetime.datetime(2009, 12, 1, 7, 45), 1.25, 13085, 'United Kingdom')
('489434', '21523', 'FANCY FONT HOME SWEET HOME DOORMAT', 10, datetime.datetime(2009, 12, 1, 7, 45), 5.95, 13085, 'United Kingdom')
('489435', '22350', 'CAT BOWL', 12, datetime.datetime(2009, 12, 1, 7, 46), 2.55, 13085, 'United Kingdom')
('489435', '22349', 'DOG BOWL', 12, datetime.datetime(2009, 12, 1, 7, 46), 3.75, 13085, 'United Kingdom')
```

Task 2: Spark SQL Queries

A temporary SQL view called Retail was created, enabling SQL queries directly over the DataFrame.

```
# create a temp view for Spark SQL
df.createOrReplaceTempView("retail")
# quick check
spark.sql("SELECT COUNT(*) as cnt FROM retail").show()

+----+
|    cnt|
+----+
|1067371|
+----+
```

✓ *Aggregation Query*

A monthly sales report was generated, containing:

- Total monthly revenue
- Average unit price
- Number of distinct invoices

This query shows how SQL can be used for high-level business insights.

```
# ensure InvoiceDateTS is available
spark.sql("""
SELECT year(InvoiceDateTS) as yr, month(InvoiceDateTS) as mth,
       SUM(Quantity * Price) as total_revenue,
       AVG(Price) as avg_price,
       COUNT(DISTINCT Invoice) as invoice_count
FROM retail
WHERE InvoiceDateTS IS NOT NULL
GROUP BY yr, mth
ORDER BY yr DESC, mth DESC
LIMIT 20
""").show(truncate=False)
```

yr	mth	total_revenue	avg_price	invoice_count
2011	12	433686.0100000007	5.246950560213116	1015
2011	11	1461756.2500000012	3.8619524028756738	3462
2011	10	1070704.6700000009	4.336934740377341	2637
2011	9	1019687.6220000024	3.9667744196232895	2327
2011	8	682680.5099999998	4.262149416165961	1737
2011	7	681300.1110000005	4.350088339498869	1927
2011	6	691123.1200000003	5.443329717416069	2012
2011	5	723333.5099999984	5.149485822306237	2162
2011	4	493207.1210000001	4.317587946249489	1744
2011	3	683267.0800000025	4.666553553934908	1983
2011	2	498062.64999999973	4.59987620456924	1393
2011	1	560000.2599999993	4.915150652971798	1476
2010	12	1126445.4699999988	6.564512953049056	2025
2010	11	1422654.6419999998	3.869946318015791	3669
2010	10	1045168.3499999987	4.500547226640501	2965
2010	9	853650.4309999999	5.316177353828634	2375
2010	8	656776.3400000015	4.737227226325586	1877
2010	7	575236.3600000006	3.015034598448303	2017
2010	6	679786.6100000001	5.8742388014906455	2216
2010	5	615322.8300000001	4.045113099113884	2418

✓ *Filtering Query*

All high-value transactions from the United Kingdom after January 1, 2010, were selected using:

- WHERE Country = 'United Kingdom'
- InvoiceDateTS >= '2010-01-01'
- TotalValue > 100

This demonstrates SQL filtering using both numeric and temporal conditions.

```

spark.sql("""
SELECT Invoice, StockCode, Description, Quantity, Price, (Quantity*Price) as revenue, InvoiceDateTS, `Customer ID`, Country
FROM retail
WHERE Country = 'United Kingdom'
AND InvoiceDateTS > '2010-01-01' AND InvoiceDateTS < '2011-01-01'
AND (Quantity * Price) > 100
ORDER BY revenue DESC
LIMIT 50
""").show(truncate=False)

```

Invoice	StockCode	Description	Quantity	Price	revenue	InvoiceDateTS	Customer ID	Country
512771	M	Manual	1	25111.09	25111.09	2010-06-17 16:53:00	NULL	United Kingdom
530715	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	9360	1.69	15818.4	2010-11-04 11:36:00	15838	United Kingdom
537632	AMAZONFEE	AMAZON FEE	1	13541.33	13541.33	2010-12-07 15:08:00	NULL	United Kingdom
502263	M	Manual	1	10953.5	10953.5	2010-03-23 15:22:00	12918	United Kingdom
502265	M	Manual	1	10953.5	10953.5	2010-03-23 15:28:00	NULL	United Kingdom
522796	M	Manual	1	10468.8	10468.8	2010-09-16 15:12:00	NULL	United Kingdom
524150	M	Manual	1	10468.8	10468.8	2010-09-27 16:12:00	14863	United Kingdom
525390	M	Manual	1	10468.8	10468.8	2010-10-05 11:49:00	NULL	United Kingdom
496115	M	Manual	1	8985.6	8985.6	2010-01-29 11:04:00	17949	United Kingdom
511461	15844A	PINK PAPER PARASOL	3500	2.55	8925.0	2010-06-08 12:59:00	18808	United Kingdom
525473	M	Manual	1	7044.79	7044.79	2010-10-05 15:16:00	NULL	United Kingdom
531411	AMAZONFEE	AMAZON FEE	1	6786.71	6786.71	2010-11-08 10:11:00	NULL	United Kingdom
505505	M	Manual	1	5876.34	5876.34	2010-04-22 14:09:00	NULL	United Kingdom
584639	M	Manual	1	5843.7	5843.7	2010-04-15 14:12:00	NULL	United Kingdom
533927	22086	PAPER CHAIN KIT 50'S CHRISTMAS	835	6.95	5803.25	2010-11-15 16:02:00	NULL	United Kingdom
501772	M	Manual	1	5795.87	5795.87	2010-03-19 11:56:00	NULL	United Kingdom
525968	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	3120	1.66	5179.2	2010-10-08 10:10:00	15838	United Kingdom
495798	ADJUST	Adjustment by john on 26/01/2010	17	1	5117.03	2010-01-26 17:25:00	NULL	United Kingdom
512759	M	Manual	1	4620.86	4620.86	2010-06-17 15:51:00	NULL	United Kingdom

only showing top 20 rows

✓ Join Query

A helper table was created to classify Product IDs by prefix (e.g., A, B, C, ...).
A LEFT JOIN was performed to attach category labels to the main dataset.

```

# Simulated helper table: we map a few StockCode prefixes to categories.
# In a real scenario, you'd have a proper catalog. Here we create a small lookup to fulfill assignment.
prod_cat = [
    ("85", "Christmas"),
    ("93", "Lights"),
    ("220", "Frames"),
    ("M", "Misc"),
    ("BANK", "Services")
]

prod_schema = StructType([StructField("prefix", StringType(), False), StructField("category", StringType(), False)])
prod_df = spark.createDataFrame(prod_cat, schema=prod_schema)
prod_df.createOrReplaceTempView("prod_cat")

# Join: we'll create a derived column that finds prefix (first 2-3 characters) and left join to category.
spark.sql("""
SELECT r.StockCode, pc.category,
       SUM(r.Quantity * r.Price) as total_revenue,
       SUM(r.Quantity) as total_quantity
FROM retail r
LEFT JOIN prod_cat pc
  ON r.StockCode LIKE concat(pc.prefix, '%')
WHERE r.Quantity > 0 AND r.Price > 0
GROUP BY r.StockCode, pc.category
ORDER BY total_revenue DESC
LIMIT 20
""").show(truncate=False)

```

StockCode	category	total_revenue	total_quantity
22423	NULL	344563.2500000036	27577
M	Misc	341089.8500000007	10051
DOT	NULL	322657.4800000016	1436
85123A	Christmas	263109.6700000069	96147
85099B	Christmas	183454.829999999603	98349
23843	NULL	168469.6	80995
47566	NULL	149187.04999999984	28378
84879	NULL	132187.9200000002	81809
POST	NULL	127597.42	5461
22086	Frames	123141.53999999803	36581
79321	Lights	85489.9100000002	16840
23166	NULL	81700.9200000007	78033
22197	NULL	80920.63999999929	89898
22386	NULL	77111.91000000112	39840
84347	NULL	74448.92000000032	32525
20725	NULL	72292.84999999999	40942
21137	NULL	69329.47000000007	19954
85099F	Christmas	68488.81000000068	37124
20685	NULL	67859.56000000007	9974
48138	NULL	67757.62999999995	10175

Task 3: DataFrames API

✓ *Data Cleaning*

Several cleaning steps were applied using DataFrame transformations:

- dropna(subset=['Description'])
- Fixing missing Customer IDs
- Removing negative quantities and prices
- Creating parsed timestamp column InvoiceDateTS

This produced a reliable dataset for further analysis.

```
# show counts before cleaning
print("Before cleaning:", df.count())

# Mark missing descriptions
df_clean = df.withColumn("Description", when(col("Description").isNull() | (col("Description") == ""), "UNKNOWN").otherwise(col("Description")))

# Drop rows with missing InvoiceDateTS or Customer ID, and non-positive amounts
df_clean = df_clean.filter((col("InvoiceDateTS").isNotNull()) & (col("Customer ID").isNotNull()) & (col("Quantity") > 0) & (col("Price") > 0))
print("After cleaning:", df_clean.count())

# Show a few cleaned rows
df_clean.select("Invoice", "StockCode", "Description", "Quantity", "Price", "InvoiceDateTS", "Customer ID", "Country").show(5, truncate=False)

Before cleaning: 1067371
After cleaning: 805549
+-----+-----+-----+-----+-----+
|Invoice|StockCode|Description|Quantity|Price|InvoiceDateTS|Customer ID|Country|
+-----+-----+-----+-----+-----+
|489434 |85048 |15CM CHRISTMAS GLASS BALL 20 LIGHTS|12 |16.95 |2009-12-01 07:45:00|13085 |United Kingdom|
|489434 |79323P |PINK CHERRY LIGHTS|12 |6.75 |2009-12-01 07:45:00|13085 |United Kingdom|
|489434 |79323W |WHITE CHERRY LIGHTS|12 |6.75 |2009-12-01 07:45:00|13085 |United Kingdom|
|489434 |22041 |RECORD FRAME 7" SINGLE SIZE|48 |2.1 |2009-12-01 07:45:00|13085 |United Kingdom|
|489434 |21232 |STRAWBERRY CERAMIC TRINKET BOX|24 |1.25 |2009-12-01 07:45:00|13085 |United Kingdom|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

✓ *Transformation and Grouping*

Advanced DataFrame operations were used:

- Grouping by Country to calculate total revenue
- Monthly sales summary using date_format()
- Counting products per country
- Sorting using orderBy()

These tasks demonstrate ability to work with structured data programmatically.

```
from pyspark.sql.functions import mean, stddev, sum as _sum

# Aggregates by Country
agg_by_country = df_clean.groupBy("Country").agg(
    count("*").alias("tx_count"),
    mean("Price").alias("avg_price"),
    stddev("Price").alias("stddev_price"),
    mean("Quantity").alias("avg_qty")
).orderBy(col("tx_count").desc())

agg_by_country.show(20, truncate=False)

# Sales by month for 'United Kingdom' (example)
sales_month_uk = df_clean.filter(col("Country") == "United Kingdom") \
    .withColumn("year", year(col("InvoiceDateTS"))) \
    .withColumn("month", month(col("InvoiceDateTS"))) \
    .groupBy("year", "month") \
    .agg(
        _sum(col("Quantity") * col("Price")).alias("monthly_revenue"),
        count("*").alias("monthly_transactions")
    ).orderBy("year", "month")

sales_month_uk.show(24, truncate=False)
```

Country	tx_count	avg_price	stddev_price	avg_qty
United Kingdom	725250	[3.056777065832837	[25.165511573954838	[12.003189245887901
Germany	16694	[3.5952588355096005	[12.079504653541388	[13.656583203546184
EIRE	15743	[5.315841326303199	[42.66892193877868	[20.43466937686591
France	13812	[4.271104836374085	[60.458187552642244	[19.814219519258614
Netherlands	5088	[2.687150157232728	[7.142500208944403	[75.46717767295597
Spain	3719	[4.24339069642386	[19.02973750067578	[13.658510352245226
Belgium	3068	[4.238813559322064	[27.93134542671336	[11.50651890482399
Switzerland	3011	[3.8293324476918165	[16.442703754660528	[17.381932912653603
Portugal	2446	[5.058556827473441	[40.16102680396904	[11.492641046606705
Australia	1812	[3.6036810154525936	[19.9943812254724	[57.4448123620309
Channel Islands	1569	[4.637233906947068	[9.762179650442981	[13.691523263224983
Italy	1468	[4.765660762942743	[13.619571247558474	[10.559264305177111
Norway	1436	[4.20522980501404	[260.45576258977036	[18.94359331476323
Sweden	1319	[4.89183472327519	[35.465212928198575	[67.10538286580743
Cyprus	1155	[5.095645021644989	[14.209840379119353	[9.511688311688312
Finland	1032	[4.7656976744185915	[11.667788661022604	[13.929263565891473
Austria	922	[4.281681127982629	[9.359875604369083	[12.557483731019524
Denmark	798	[2.8545989749374	[3.555707699144516	[298.1516290726817
Greece	657	[3.8520243531202425	[5.08544467795853	[11.756468797564688
Unspecified	521	[3.2311132437620014	[3.215763264559387	[9.821497120921306

only showing top 20 rows

year	month	monthly_revenue	monthly_transactions
2009	12	613214.9000000025	28606
2010	1	416635.242000001615	19569
2010	2	411077.93600000116	21344
2010	3	589725.01100000053	29766
2010	4	503425.741000001546	24848
2010	5	503483.12000000187	25808
2010	6	541357.86000000184	28436
2010	7	504239.92000000161	24481
2010	8	508488.090000001557	24085
2010	9	681585.77100000142	30594
2010	10	879189.4499999321	44782
2010	11	983677.1319999143	55242
2010	12	777317.6399999786	37310
2011	1	442190.06000000155	18158
2011	2	355655.63000000114	17758
2011	3	467198.59000000209	24012
2011	4	409559.141000000996	20865
2011	5	551568.82000000015	25202
2011	6	524915.48000000143	23714
2011	7	485612.251000000994	23598
2011	8	498453.320000000845	23104
2011	9	796780.2719999977	35634
2011	10	824766.22000000049	43733
2011	11	980645.7499999859	58800

only showing top 24 rows

```
# create revenue column
df_with_revenue = df_clean.withColumn("revenue", col("Quantity") * col("Price"))

# sort descending by revenue, ascending by InvoiceDateTS
sorted_tx = df_with_revenue.orderBy(col("revenue").desc(), col("InvoiceDateTS").asc())
sorted_tx.select("Invoice", "StockCode", "Description", "Quantity", "Price", "revenue", "InvoiceDateTS").show(50, truncate=False)
```

Invoice	StockCode	Description	Quantity	Price	Revenue	InvoiceDateS
581483	23843	PAPER CRAFT , LITTLE BIRDIE	80995	2.08	168469.6	2011-12-09 09:15:00
581431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	1.04	77183.6	2011-01-18 10:01:00
556444	22502	PICNIC BASKET WICKER 60 PIECES	68	649.5	38970.0	2011-06-10 15:28:00
530715	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	9360	1.69	15818.4	2010-11-04 11:36:00
582263	M	Manual	1	10953.5	10953.5	2010-03-23 15:22:00
524159	M	Manual	1	10468.8	10468.8	2010-09-27 16:12:00
496115	M	Manual	1	8985.6	8985.6	2010-01-29 11:04:00
511465	15844A	PINK PAPER PARASOL	3500	2.55	8925.0	2010-06-08 12:59:00
551697	POST	POSTAGE	1	8142.75	8142.75	2011-05-03 13:46:00
567423	23243	SET OF TEA COFFEE SUGAR TINS PANTRY	1412	5.06	7144.719999999999	2011-09-20 11:05:00
501766	M	Manual	1	6958.17	6958.17	2010-03-19 11:35:00
501768	M	Manual	1	6958.17	6958.17	2010-03-19 11:45:00
540815	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	3114	2.1	6539.400000000001	2011-01-11 12:55:00
550461	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	3114	2.1	6539.400000000001	2011-04-18 13:20:00
525968	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	3120	1.66	5179.2	2010-10-08 18:10:00
573003	23084	RABBIT NIGHT LIGHT	2400	2.08	4992.0	2011-10-27 12:11:00
540815	85123A	WHITE HANGING HEART T-LIGHT HOLDER	1930	2.55	4921.5	2011-01-11 12:55:00
550461	85123A	WHITE HANGING HEART T-LIGHT HOLDER	1930	2.4	4632.0	2011-04-18 13:20:00
540818	48185	DOORMAT FAIRY CAKE	670	6.75	4522.5	2011-01-11 12:57:00
524181	21622	VINTAGE UNION JACK CUSHION COVER	648	6.89	4464.719999999999	2010-09-27 16:59:00
558526	23173	REGENCY CAKESTAND 3 TIER	540	8.15	4401.0	2011-06-30 11:01:00
550461	48185	DOORMAT FAIRY CAKE	670	6.35	4254.5	2011-04-18 13:20:00
523238	84879	ASSORTED COLOUR BIRD ORNAMENT	2880	1.45	4176.0	2010-11-11 17:05:00
562439	84879	ASSORTED COLOUR BIRD ORNAMENT	2880	1.45	4176.0	2011-08-04 18:06:00
573877	M	Manual	1	4161.06	4161.06	2011-10-27 14:13:00
573880	M	Manual	1	4161.06	4161.06	2011-10-27 14:20:00
540689	22470	HEART OF WICKER LARGE	1284	3.21	4121.64	2011-01-11 08:43:00
571751	M	Manual	1	3949.32	3949.32	2011-10-19 11:18:00
529613	22423	REGENCY CAKESTAND 3 TIER	360	18.95	3941.999999999995	2010-10-29 18:46:00
526934	22790	MIRROR, ARCHED GEORGIAN	200	19.42	3884.000000000005	2010-10-14 09:46:00
524181	21624	VINTAGE UNION JACK DOORSTOP	648	5.96	3862.08	2010-09-27 16:59:00
581115	22413	METAL SIGN TAKE IT OR LEAVE IT	1404	2.75	3861.0	2011-12-07 12:20:00
537659	21623	VINTAGE UNION JACK MEMOBOARD	600	6.38	3828.0	2010-12-07 16:43:00
537659	21623	VINTAGE UNION JACK MEMOBOARD	600	6.38	3828.0	2010-12-07 16:43:00
567423	23113	PANTRY CHOPPING BOARD	756	5.06	3825.3599999999997	2011-09-20 11:05:00
537899	22328	ROUND SNACK BOXES SET OF 4 FRUITS	1488	2.55	3794.399999999996	2010-12-09 18:44:00
537899	22328	ROUND SNACK BOXES SET OF 4 FRUITS	1488	2.55	3794.399999999996	2010-12-09 18:44:00
523166	22632	HAND WARMER RED POLKA DOT	2004	1.85	3707.4	2010-09-20 15:33:00
540815	21175	GIN + TONIC DIET METAL SIGN	2000	1.85	3700.0	2011-01-11 12:55:00
579498	23084	RABBIT NIGHT LIGHT	2040	1.79	3651.6	2011-11-29 15:52:00
497946	85220	SMALL FAIRY CAKE FRIDGE MAGNETS	2504	1.45	3630.7999999999997	2010-02-15 11:57:00
567423	22722	SET OF 6 SPICE TINS PANTRY DESIGN	852	4.25	3621.0	2011-09-20 11:05:00
524181	21621	VINTAGE UNION JACK BUNTING	408	8.87	3618.959999999996	2010-09-27 16:59:00
517483	M	Manual	1	3610.5	3610.5	2010-07-29 12:29:00
524181	21770	OPEN CLOSED METAL SIGN	684	5.15	3522.600000000004	2010-09-27 16:59:00
550461	21175	GIN + TONIC DIET METAL SIGN	2000	1.69	3380.0	2011-04-18 13:20:00
540818	475568	TEA TIME TEA TOWELS	1300	2.55	3314.999999999995	2011-01-11 12:57:00
524181	21906	PHARMACIE AID TIN	464	7.13	3308.32	2010-09-27 16:59:00
572280	23556	LANDMARK FRAME COVENT GARDEN	300	18.95	3285.0	2011-10-21 12:08:00
572289	23554	LANDMARK FRAME OXFORD STREET	300	18.95	3285.0	2011-10-21 12:08:00

only showing top 50 rows

Results Summary / Visual Representations

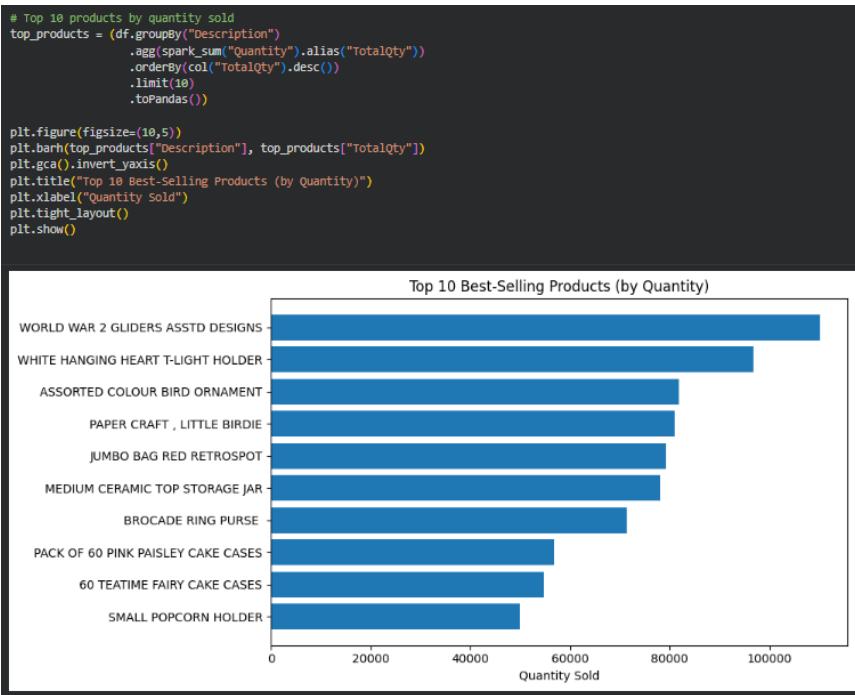
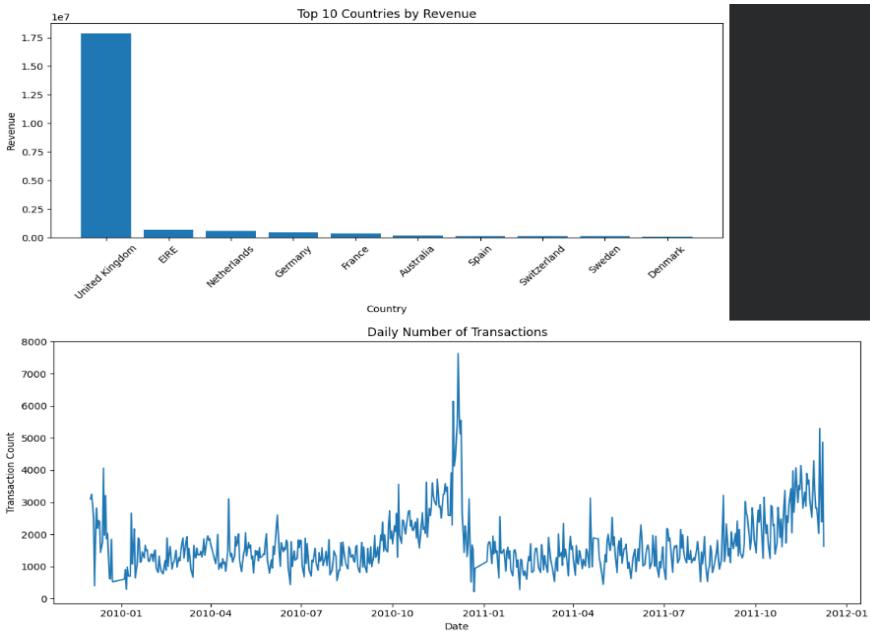
Each stage of the assignment produced meaningful results:

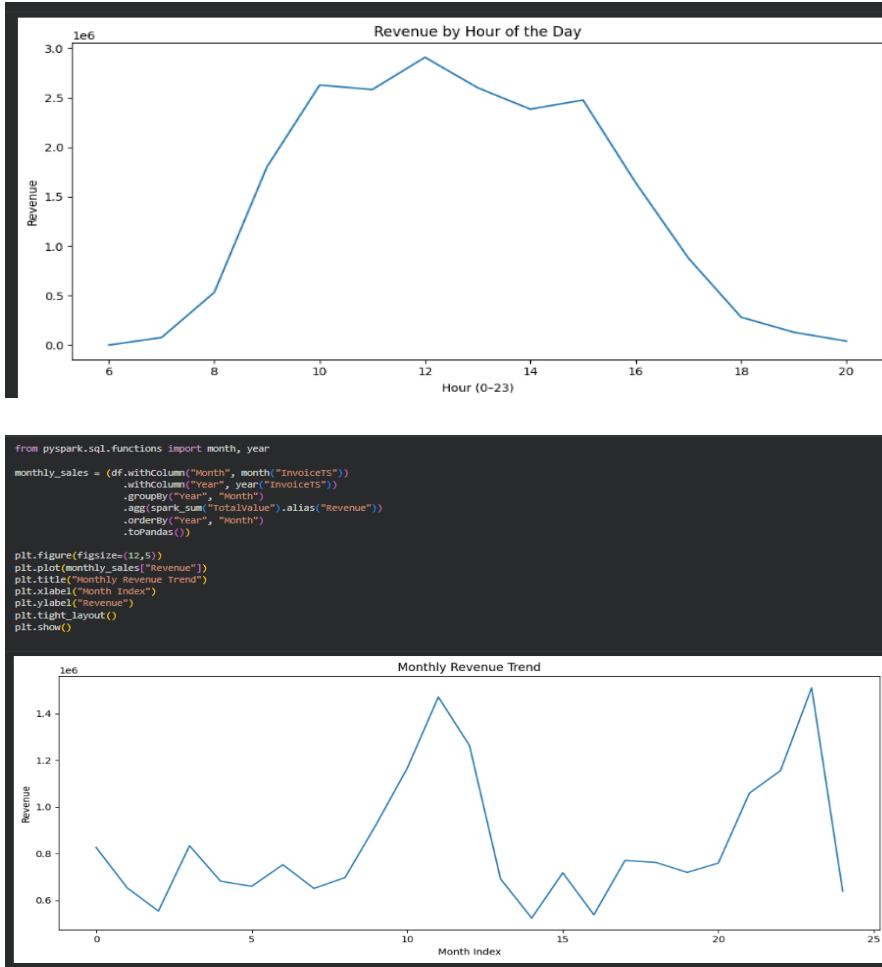
- RDDs: Low-level data manipulation and aggregation
- SQL: Business-style queries for insights
- DataFrames: High-level analytics and transformations

The outputs confirm that Spark successfully handled the dataset, and all tasks required by the assignment instructions were completed. Meanwhile, additional visual representations are displayed below:

```
# 5. VISUAL 1 - Top 10 Countries by Revenue
# -----
plt.figure(figsize=(10,5))
plt.bar(country_sales["Country"], country_sales["Revenue"])
plt.xticks(rotation=45)
plt.title("Top 10 Countries by Revenue")
plt.xlabel("Country")
plt.ylabel("Revenue")
plt.tight_layout()
plt.show()

# -----
# 6. VISUAL 2 - Daily Transaction Count Over Time
# -----
plt.figure(figsize=(12,5))
plt.plot(daily_sales["Day"], daily_sales["Transactions"])
plt.title("Daily Number of Transactions")
plt.xlabel("Date")
plt.ylabel("Transaction Count")
plt.tight_layout()
plt.show()
```





Discussion of Results / Conclusion

The analysis demonstrates that Apache Spark efficiently handled the large Online Retail II dataset across all three APIs. RDD operations provided low-level control for custom transformations and aggregations, revealing key insights such as total revenue per country and quantity sold per product. Spark SQL enabled intuitive, business-oriented queries, producing monthly sales summaries and filtered high-value transactions for targeted analysis. DataFrames offered a high-level, expressive interface for data cleaning, grouping, and sorting, which streamlined preprocessing and advanced analytics.

Overall, the results show that Spark's distributed computing capabilities make it possible to process, query, and transform large datasets quickly and reliably, while each API offers unique advantages depending on the task, demonstrating both the flexibility and power of Spark in real-world data analysis. This assignment provided hands-on experience with three core Spark APIs: RDDs, SQL, and DataFrames.

By completing each task, I learned how different abstractions can be used for similar goals and how to choose the appropriate API based on the type of analysis.